

IranWar.ai: An Open-Source Event-Level Dataset of the 2026 US–Iran Conflict

Jacob E. Thomas, MA, PhD; Aigerim Alpysbekova, MPH, PhD(c); Eugene Osei Mensah, BSN, RN;
Nigel Masara, BSc, MSc; Prakhar Sharma, PhD

Original Preprint Release Date: 05/13/2026

Dashboard: <https://iranwar.ai> | Repository: github.com/jethomasphd/WarTheater

Abstract

We present the IranWar.ai Event-Level Research Dataset, a structured, openly available record of the first 30 days of the 2026 US–Iran conflict (Operation Epic Fury, February 28–March 29, 2026). The dataset comprises 1,653 event-level observations across 48 variables, derived from 16 curated JSON data files that power the IranWar.ai public intelligence dashboard. Each observation represents a discrete, verifiable event: an airstrike on a named target, a retaliatory attack, a financial market data point, a daily casualty estimate by faction, a naval deployment, or a diplomatic development. Nine analytical domains are covered (military strikes, retaliations, financial markets, humanitarian impacts, diplomatic events, naval operations, cyber activity, general military, and contextual data) integrating open-source intelligence (OSINT) from more than 60 specific source streams, including CENTCOM, the IDF, ACLED, the Iranian Red Crescent, the IAEA, Bloomberg, and the IMO. Beyond its function as a research resource, this paper examines the dataset as an epistemic object. Assembled in part by AI research agents and carrying known imperfections, the dataset is released with full transparency about its own limitations, including interpolated values, unverified records, and probabilistic confidence ratings. We argue that this transparency constitutes a deliberate counter-technology: a structural intervention against the conditions of the information flood, which we frame as a dual-layer epistemic crisis comprising both a structural component (the volume, velocity, and synthetic origin of available data overwhelming the infrastructure available to evaluate it) and a strategic component (the deliberate production of disinformation by state and non-state actors). Radical transparency is necessary for both layers but, by itself, is sufficient only against the structural layer; the strategic layer additionally requires interpretive expertise that situates discrepancies as political facts rather than data-quality defects. We further examine the dataset's relationship to AI training pipelines, model collapse dynamics, and the degradation of provenance metadata across summarization and citation chains. To address the limits of any single research team's expertise, we set out a tentative, non-exhaustive research agenda and invite collaborators across political science, public health, economics, military studies, media studies, and information science. The live dashboard, codebook, source data, prompts, manifests, reproducible extraction script, and versioned research releases are publicly available at github.com/jethomasphd/WarTheater; scheduled research releases are monthly, with interim correction releases when warranted.

Keywords: conflict data, open-source intelligence, event-level dataset, US–Iran, OSINT, information flood, synthetic data, epistemic infrastructure, model collapse, reproducible research, civilian harm exposure

1 Introduction

The US–Iran conflict that began on February 28, 2026 (designated Operation Epic Fury by the US Department of Defense) represents the largest American military engagement since the 2003 invasion of Iraq. Within its first 30 days, the conflict produced over 30,500 US and Israeli strike targets across 26 of Iran's 31 provinces, an effective closure of the Strait of Hormuz affecting 20% of global oil supply, a 58% surge in Brent crude oil prices, the opening of secondary fronts in Lebanon and Yemen, and casualty estimates in the thousands across multiple countries.

The pace and complexity of this conflict create an urgent need for structured, research-ready data. Existing sources (e.g., news reports, think-tank analyses, government press releases) are fragmented across organizations, inconsistent in format, and difficult to analyze quantitatively. These problems are not abstract. CENTCOM and Iranian state media have repeatedly described the same strike with incompatible target lists; ACLED, the Iranian Red Crescent, and HRANA have reported casualty figures that diverge by an order of magnitude on the same dates; geographic information appears in some sources as place names, in others as administrative units, and in still others as coordinates with varying precision; timestamps are reported in at least four time zones (Iran/IRST, Gulf/GST, Israel/IST, US/ET) without normalization. A researcher who wants to estimate strike intensity over time, or to test whether oil-price moves track specific targeting patterns, must first reconcile these sources by hand which does not scale to a conflict producing thousands of events per month. Real-time tracking dashboards provide situational awareness but are designed for visualization, not statistical analysis. The challenges of assembling reliable event-level data from media reports are well documented (Weidmann, 2015); in a conflict of this scale and tempo, those challenges are compounded (Hegre, Croicu, Eck, & Höglbladh, 2020; Miller, Kishi, Raleigh, & Dowd, 2022).

This conflict is also distinctive in a way that matters for the data infrastructure built around it. Three features make the information environment unusually hostile to ordinary curation. First, the conflict is occurring during the first sustained interstate war of the generative-AI era: synthetic video, synthetic audio, and AI-generated still imagery are entering the open-source pipeline at industrial scale. Within the first two weeks of fighting, Cyabra reported that a coordinated network of inauthentic accounts distributing AI-generated war footage generated more than 145 million views across major platforms; Murray and Linvill (2026) separately documented a multilingual, IRGC-affiliated influence operation across X, Instagram, and Bluesky (Cyabra, 2026; Murray & Linvill, 2026). Second, the conflict is multi-theater and multi-actor in a way that fragments reporting across at least seven national press systems and multiple time zones. Third, the political stakes of the conflict, including the closure of the Strait of Hormuz and the involvement of nuclear-related infrastructure, generate strategic incentives to misreport that are stronger than in most prior conflicts whose data has eventually entered the academic record (Gibilisco & Steinberg, 2023).

The combined effect is a research environment in which the time available to verify any given claim is shorter, and the cost of misclassification higher, than for the conflicts on which existing event datasets were originally built.

The IranWar.ai project responds to this need. Originally built as a public-facing OSINT intelligence dashboard, the project has made its underlying data and code available from the opening days of the conflict, while the dashboard itself translates those data into interactive maps, charts, timelines, and briefings for public use. This paper describes both the dashboard architecture and the transformation of its structured JSON files into a research-ready, event-level CSV dataset suitable for quantitative analysis. The dashboard is designed for daily updating; the research dataset is designed for scheduled, versioned monthly releases, with interim correction releases when major errors, source updates, or analytically important reclassifications warrant earlier publication. Compared to existing event datasets (principally UCDP-GED and ACLED) our contribution is distinguished by (i) a near-real-time, open-source data pipeline with a documented two-phase production protocol that other research teams can audit and replicate; (ii) explicit, record-level confidence scoring, source attribution, and AI-assist disclosure; and (iii) the integration of military, financial, humanitarian, and diplomatic events in a single event-level structure rather than as separate files. We discuss the trade-offs of this approach, particularly the exposure to interpretive error that attends real-time work, throughout the paper.

But a dataset is never merely a dataset. Every structured record that enters the information ecosystem carries implicit claims about what happened, when, where, and to whom: claims that downstream consumers may accept, contest, aggregate, summarize, or train machine learning models on. The present dataset is no exception. It was assembled in large part by AI research agents processing open-source reporting. It contains interpolated values, unverified records, and confidence ratings that reflect probability, not certainty. Some of its content may be wrong. We know this, and we are releasing it anyway. Not despite these imperfections, but in part because of them.

We do so because the alternative (i.e., withholding imperfect data until it achieves a standard of verification that no conflict dataset has ever met in real time) concedes the information environment to sources that carry no provenance markers at all. The condition we are describing is what we term the *information flood*¹: a dual-layer epistemic crisis in which (a) the volume, velocity, and synthetic origin of available data overwhelm the infrastructure available to evaluate it, and (b) state and non-state actors deliberately produce disinformation at scale to advance strategic interests. The two layers interact: structural overload makes adversarial production cheaper, and adversarial production accelerates structural overload. In this environment, transparency about uncertainty is not a weakness. It is a *counter-technology*²: a deliberate structural choice to restore the conditions under which data can be meaningfully assessed. As we

develop in Section 6, transparency is necessary against both layers of the flood, but sufficient only against the first; the second additionally requires interpretive expertise that this paper attempts to model.

We therefore release this dataset with four goals: (1) to provide researchers with immediate access to structured conflict data without the lag typical of academic datasets; (2) to establish an open, versioned, and reproducible data infrastructure with monthly updates for the duration of the conflict; (3) to invite collaborative contributions from researchers across disciplines; and (4) to examine what happens when a dataset that is honest about its own imperfections enters an information ecosystem that rarely demands such honesty, and to ask whether that honesty changes anything.

1.1 Related Work

Event-level conflict datasets have a strong tradition in political science and peace research. The Uppsala Conflict Data Program (UCDP) Georeferenced Event Dataset (Sundberg & Melander, 2013) and the Armed Conflict Location & Event Data Project (ACLED; Raleigh et al., 2010) provide global coverage of political violence events, but with reporting lags of months to years. The UCDP Candidate Events Dataset (Hegre, Croicu, Eck, & Högladh, 2020) reduces this lag to roughly one month for organized-violence events. The Militarized Interstate Dispute (MID) dataset (Palmer et al., 2022) captures interstate conflicts at lower temporal resolution, building on the broader Armed Conflict dataset tradition (Gleditsch et al., 2002). Comparative evaluations of these resources (Eck, 2012) have highlighted meaningful differences in coding rules, coverage, and quality control; differences that have practical consequences for downstream analysis. More recent work argues that bias in conflict data is shaped by collection design, source breadth, and the trade-off between comprehensiveness and stability, and calls for an explicit research agenda on these biases (Miller, Kishi, Raleigh, & Dowd, 2022). For economic dimensions of conflict, researchers typically rely on FRED, Bloomberg, or custom data collection.

Our dataset differs from these resources in four respects. First, the underlying data have been available as open-source dashboard data from the opening days of the conflict, while the paper described here is one versioned research release extracted from that continuously auditable public record. In other words, the research releases have a monthly cadence, but the source data and dashboard are designed for higher-frequency public availability. Second, the release is near-real-time relative to retrospective conflict datasets, with scheduled monthly snapshots and interim correction releases when warranted. Third, it integrates military, financial, humanitarian, and diplomatic data in a single event-level structure, enabling cross-domain analysis that existing event datasets do not support directly. We see this as one of the dataset's most important contributions, because the most analytically interesting questions about this conflict sit at the intersections of

these domains. Concretely, the cross-domain structure enables research questions such as the following:

- **Military–Financial:** do strikes near the Strait of Hormuz produce immediate spikes in Brent crude or in tanker insurance, and on what timescale do those spikes decay?
- **Military–Humanitarian:** which categories of targets (energy, transport, healthcare, governance) are most strongly associated with subsequent civilian displacement and casualty escalation?
- **Military–Diplomatic:** do strike clusters precede, follow, or co-occur with diplomatic openings, and does the directionality differ across actors (US, Israel, Iran, third parties)?
- **Financial–Diplomatic:** do new sanctions correlate with measurable currency or capital-flight signals, and do those signals precede or follow public diplomatic announcements?
- **Humanitarian–Diplomatic:** does rising civilian harm or displacement translate into increased international aid commitments, increased diplomatic intervention, or neither?

Third, our dataset is fully reproducible: the extraction script, source data, codebook, prompts, update manifests, and database snapshots are versioned in a public Git repository. A fourth distinction, explored in Section 6, is that this paper explicitly examines the dataset's own epistemic status: its relationship to truth, error, and the information environment it enters.

2 Data Collection

2.1 The IranWar.ai Dashboard

The IranWar.ai dashboard is a static-site intelligence platform built with HTML, CSS, and JavaScript (no backend framework), hosted on Cloudflare Pages, and powered entirely by structured JSON files stored in a public GitHub repository. The dashboard has been updated daily since Day 1 of the conflict using a two-phase protocol:

Phase 1 (Deep Research). An AI research agent powered by Anthropic Claude Opus 4.6 processes the latest reporting from curated source streams (CENTCOM press releases, IDF statements, ACLED event data, Reuters/AP wire reports, government ministry statements, satellite imagery analysis, and financial data feeds) and produces an Update Manifest: a structured document specifying every data change for the day. The manifest is also the first place where the strategic-reporting incentives identified by Gibilisco and Steinberg (2023) are explicitly handled: when sources disagree systematically (e.g., on casualty counts), the manifest records both figures and the source of each rather than averaging them away.

Phase 2 (Code Execution). A coding agent powered by Anthropic Claude Opus 4.6 applies the manifest to the JSON data files, validates all modifications, and commits to the repository. Changes deploy automatically to the live dashboard via Cloudflare Pages.

This protocol ensures that every data point is traceable to a specific source record or source category when such attribution is available, and the full history of changes is preserved in version control. It also means that the research agent's interpretive judgments (i.e., which sources to trust, how to reconcile conflicting reports, what confidence level to assign) are embedded in the data. The prompts used by both agents, the repository history, and the database snapshots form the project's provenance trail. We return to the implications of this in Section 6.

The research dataset is extracted from these JSON files at monthly intervals. Each monthly release is versioned (v1.0, v1.1, etc.), with a snapshot of the source data preserved at the time of extraction. This monthly cadence ensures that the dataset grows as a longitudinal record: researchers working with v1.0 can reproduce their analyses exactly, while later versions incorporate new events, corrections, and reclassifications.

2.2 Source Attribution

The dashboard draws from more than 60 specific source streams. Table 1 summarizes the major source categories.

Table 1. Source categories for the IranWar.ai dataset.

Source Category	Examples
US Government	CENTCOM, DoD, Pentagon, White House, CISA, State Dept.
Israeli Government	IDF, Magen David Adom
Iranian Sources	Red Crescent, Health Ministry, IRNA, HRANA
International Organizations	IAEA, WHO, OCHA, UNHCR, IMO, IEA
Conflict Monitoring	ACLED, CTP-ISW, Bellingcat, BBC Verify
Financial Data	Bloomberg, ICE Futures, NYMEX, AAA, EIA, FRED
Military Tracking	USNI Fleet Tracker, CRS Reports, DVIDSHUB
Wire Services	Reuters, AP, AFP, Al Jazeera, CNN, NYT

Source attribution is preserved at the record level in most data files and propagated to the research dataset. This record-level attribution is a deliberate design choice informed by recent

work on dataset documentation standards (Gebru et al., 2021), which argues that transparency about data provenance is a precondition for responsible downstream use.

3 Dataset Construction

3.1 Design Principles

The dataset was built around a few core commitments that shaped every design decision. The most important was event-level granularity: each row represents a single observable event, not an aggregate. This matters because aggregation is irreversible: a researcher who needs finer resolution cannot recover it from summary data, but anyone who needs summaries can always roll up from events. We wanted to give analysts the most disaggregated foundation we could.

The second commitment was cross-domain integration. It would have been easier to release separate files for military data, financial data, and humanitarian data. We chose not to, because the most interesting questions in this conflict sit at the intersections: how did oil markets respond to specific strike patterns? Did diplomatic events precede or follow escalations? A fragmented release would have made those questions harder to ask.

Third, we preserved everything. No fields from the source JSON files were discarded. Data that didn't fit the common schema was retained in domain-prefixed extra columns rather than silently dropped. We made this choice partly for scientific reproducibility and partly out of humility: we don't always know in advance which fields will matter.

Finally, the extraction script (*build_dataset.py*) is included in the repository and produces identical output from identical inputs. Every monthly release will include both the script and a snapshot of the source data it was run on.

3.2 Schema Design

Every row contains nine required columns (Table 2). These are supplemented by 17 domain-specific columns (location, actors, weapons, casualties, financial metrics, infrastructure type) and 22 additional columns that preserve domain-specific detail. All domain-specific columns are NULL when not applicable to a given event type.

Table 2. Required columns in the event-level schema.

Column	Description
event_id	Unique sequential identifier (EVT-0001, EVT-0002, ...)
date	ISO 8601 date (YYYY-MM-DD)

Column	Description
datetime_utc	Full timestamp when available
day_of_conflict	Integer (Day 0 = pre-war baseline, Day 1 = first strikes)
event_domain	Primary classification (9 categories; see Table 3)
event_type	Subcategory within domain (25 unique values)
event_description	Free-text description of the event
source_file	JSON file from which the row was extracted
source_record_id	Original identifier from source data

3.3 Event Domain Taxonomy

The nine event domains and their observation counts are shown in Table 3.

Table 3. Event domain taxonomy with observation counts (v1.0).

Domain	N	Description
STRIKE	641	US/Israeli offensive military strikes on Iranian territory
RETALIATION	307	Iranian, Hezbollah, Houthi, and proxy counter-attacks
FINANCIAL	301	Oil prices, market indices, war costs, tanker transits
HUMANITARIAN	173	Daily casualty reports by faction, infrastructure damage
OTHER	69	Daily briefing headlines, aggregate snapshots
DIPLOMATIC	66	Political statements, ceasefire negotiations
MILITARY	64	General military events, US military base reference data
NAVAL	28	Fleet deployments, Strait of Hormuz events
CYBER	4	Offensive and defensive cyber operations

3.4 Explosion Logic

Two source files (*strikes-iran.json* and *strikes-retaliation.json*) use a compact representation designed for the dashboard: a single record describes a location that was active across multiple days. To convert this into event-level rows, we apply what we call "explosion logic."

For strikes on Iran, each of 61 location records contains arrays of both targets (specific installations) and *active days* (the days that the location was struck). Taking the Cartesian product of these arrays (i.e., one row for each target on each active day) yields 620 rows. For retaliation strikes, each of 72 location records has an *active_days* array but no target sub-array; explosion yields one row per active day, for 237 rows.

This granularity is what makes temporal analysis possible. Without it, a researcher couldn't track how strike tempo changed over the 30-day window, or identify shifts in the target category. Future monthly releases will extend the *active_days* arrays as new strike activity is recorded, and row counts will grow accordingly.

3.5 Data Confidence Scoring

Every row receives a *data_confidence* rating:

HIGH (76% of rows). Extracted directly from structured data fields with authoritative source attribution (CENTCOM, DoD, IDF, IAEA, verified financial feeds).

MEDIUM (20%). Inferred from secondary sources (ACLED, satellite imagery and interpolated financial data).

LOW (4%). From unverified reports, Wikipedia-sourced entries, or records explicitly flagged as unconfirmed in the source data.

These ratings are carried forward in every monthly release. Records initially rated LOW or MEDIUM may be upgraded in subsequent versions as additional reporting confirms or refutes them; such reclassifications will be logged in the release changelog.

4 Summary Statistics

4.1 Temporal Coverage

The dataset's temporal window runs from January 2, 2026 (i.e., the earliest pre-war oil price observation) through March 29, 2026, Day 30 of the conflict. The conflict event window itself begins on February 28 with the first US strikes. Each subsequent monthly release will extend this window by approximately 30 days.

4.2 Geographic Coverage

Strike and retaliation records span 13 countries or zones: Iran, Israel, Kuwait, Qatar, UAE, Bahrain, Iraq, Saudi Arabia, Lebanon, Jordan, Cyprus, Oman, and international waters in the Strait of Hormuz region. Within Iran, the 61 distinct strike locations documented in this version span 26

of the country's 31 provinces: a geographic footprint significantly broader than many early reporting accounts suggested.

4.3 Column Completeness

The nine core columns (*event_id*, *date*, *day_of_conflict*, *event_domain*, *event_type*, *event_description*, *source_file*, and *data_confidence*) are 100% complete across all rows. Geographic columns are populated for roughly 53–55% of rows; geographic coordinates are simply not meaningful for financial time-series or diplomatic events, so this is by design rather than a gap. Casualty data is present for 26% of rows and financial metrics for 15%.

4.4 Source File Contribution

The *timeline-events.json* file contributes the most diverse cross-section of events (293 rows spanning all nine domains), while the combined strike files account for 857 rows (52% of the dataset by volume). This distribution reflects the dataset's core purpose: military activity is the most densely documented dimension of this conflict, and the dataset accurately reflects that asymmetry.

5 Known Limitations

No conflict dataset assembled from open-source reporting is free of error, and this one is no exception. The limitations below are documented not to discourage use but to support it: researchers who understand the dataset's constraints will use it more accurately, and will be better positioned to push back if their findings are challenged. Several of these limitations are structural features of all media-based conflict data (Weidmann, 2015; Eck, 2012; Miller et al., 2022), not deficiencies unique to this project:

1. **Casualty figures are from initial reports.** Iranian casualty data comes primarily from the Iranian Red Crescent and HRANA. These figures are difficult to verify independently due to internet blackouts and the fog of war. Discrepancies between sources (e.g., Red Crescent reporting 1,900+ killed vs. HRANA reporting 3,461) are documented in the data but not reconciled. We treat such discrepancies as politically informative signal rather than as defects to be averaged away (see Section 6.5).
2. **Daily casualty data represents estimates, not counts.** The *casualties.json* data is derived from chart-display data on the dashboard, representing the estimated daily killed per faction. These are not ground-truth body counts. Cumulative totals in the snapshot columns may not equal the sum of daily estimates.

3. **Four retaliation location records are unverified.** The *strike_verified* column flags records where open-source confirmation was not found as of the extraction date. These may be resolved in future monthly releases.
4. **Financial data has gaps on non-trading days.** Oil prices and market indices are NULL on weekends and holidays, since exchanges are closed. Market sector data is expressed as an index normalized to February 27 = 100, not as raw index values.
5. **War cost estimates are interpolated.** The cost time series is anchored to two public figures: \$11.3 billion through Day 6 (Pentagon) and \$16.5 billion through Day 12 (CSIS). Values between and beyond these anchors are estimated using a daily run rate derived from those points. Researchers should treat the interpolated values as approximations with meaningful uncertainty.
6. ***active_days* indicate location-level activity, not sortie counts.** A single active day for a given location may represent one strike, several strikes, or a sustained bombardment. The dataset does not currently resolve this distinction.
7. **The dataset cannot capture what was not reported.** Events that evaded media coverage, classified operations, and actions in areas under communications blackout are not represented. Future versions may incorporate structured civilian observation reporting (see Section 8.2) as a complementary stream, treating such reports as additional signal to be triangulated against existing sources rather than as verified ground truth.
8. **Time zone imprecision.** Timeline event timestamps reflect the time zone of the reporting source (which varies: Iran/IRST, Gulf/GST, Israel/IST, US/ET) and are not normalized to UTC.

5.1 The Politics of Event Classification

The dataset's methodological transparency, depicted in its confidence scores, source attribution, and documented extraction process, represents a significant advance over conflict datasets that present coded events as neutral facts. But there is a deeper layer of interpretive politics that transparency alone cannot resolve: the classification of events is itself a political act.

Every event in this dataset was classified across three dimensions where contestation is not incidental but structural. The first is *description*: did this event happen? The divergence between CENTCOM press releases and Iranian government statements on the same strikes illustrates that the answer is rarely unambiguous.

The second is *categorization*: what kind of event was this? Whether a strike on a power station is coded as a STRIKE or a HUMANITARIAN event (in other words, whether it is treated

as a military operation or as a civilian-infrastructure attack) depends on interpretive choices that carry political valence. The US and Iranian governments would code the same event differently. ACLED and CENTCOM would code it differently. The extraction pipeline that assembled this dataset made those calls under documented instructions designed by the authors, but the choices themselves are not politically neutral.

The third dimension is *attribution*: who was responsible? In a conflict involving US forces, Israeli forces, Iranian proxies, and independent non-state actors operating across multiple fronts, attribution is frequently contested and occasionally unknowable in real time.

This is not a criticism of the dataset's construction. It is an argument for how it should be used. Researchers applying this data to coding-sensitive analyses (i.e., anything involving attribution, target category, or civilian-versus-military distinction) should treat the coding decisions not as ground truth but as one defensible interpretation among several. Where the political stakes of classification are high, intercoder reliability checks using source-disaggregated records are advisable: coding the same event as CENTCOM would, as the Iranian Red Crescent would, and as ACLED would, and using the variance as a measure of political contestation rather than as a problem to be resolved.

This approach draws on a well-established tradition in conflict data methodology. Eck's (2012) comparative evaluation of UCDP and ACLED documents how coding rules and source selection produce systematically different event records from the same underlying conflict. Gibilisco and Steinberg's (2023) strategic-reporting model extends this insight: if governments strategically frame their conflict reporting, then any dataset that draws on government sources inherits those strategic framings to some degree. The solution is not to avoid government sources, which would make real-time conflict data collection impossible, but to treat source identity as an analytical variable rather than as a nuisance to be controlled away.

The IranWar.ai dataset is unusually well-positioned for this kind of analysis precisely because it preserves source attribution at the record level. Researchers can disaggregate by source category between US government, Israeli government, Iranian sources, international organizations, wire services, and examine how event descriptions, casualty figures, and target classifications vary across source types. That variance is not an error; it is one of the most politically informative signals in the dataset.

6 The Dataset as Epistemic Object

The preceding sections describe what the dataset contains and how it was built. This section examines what the dataset *is*: not as a technical artifact, but as an object that enters the information ecosystem and participates in the construction of knowledge about a real conflict. We argue that

this examination is not peripheral to the data contribution. It is, in a meaningful sense, the contribution's other half.

6.1 Provenance and the Interpretive Chain

Every row in this dataset passed through an interpretive chain before arriving in its present form. A journalist or government spokesperson produced a report. A wire service or monitoring organization aggregated it. An AI research agent (i.e., a large language model operating under structured instructions) read the aggregated reporting, made judgments about salience and confidence, and produced a structured Update Manifest. A coding agent applied the manifest to JSON data files. An extraction script transformed those files into the CSV rows presented here.

At no point in this chain did a human with direct observational access to events in Iran encode a data point. The dataset is, in a precise sense, a record of what was reported about what happened: filtered through multiple layers of institutional, algorithmic, and editorial interpretation. This is not unique to our dataset. It is the condition of all conflict data assembled from open sources (Weidmann, 2015). But it is rarely stated this plainly.

We state it plainly because the alternative (i.e., presenting the data as a transparent window onto events) would be a more dangerous form of imprecision than any error in our casualty figures. The flood does not lack data. It lacks provenance. The call for systematic dataset documentation (Geburu et al., 2021) reflects a growing recognition that provenance infrastructure is not optional but foundational.

6.2 Known Errors as Structural Features

Section 5 catalogs the dataset's known limitations in the conventional manner of a data descriptor. But consider what those limitations actually represent. The casualty discrepancy between the Iranian Red Crescent (1,900+) and HRANA (3,461) is not a data quality problem that will be resolved by better collection methods. It is a political fact: a direct expression of the contested relationship between a state apparatus and a diaspora human rights organization, each with different access, different incentives, and different definitions of "confirmed."

The interpolated war cost estimates are not placeholders awaiting real data. They are the best available approximation, anchored to two government figures, and they may never be superseded, because the actual cost of a military operation is itself a contested, definitionally unstable quantity that depends on what you count, over what time horizon, and for whom.

The four unverified retaliation locations are not bugs. They are the dataset's honest admission that the fog of war is not a metaphor. It is an operational condition in which some things

are reported, some are fabricated, some are real but unreported, and the boundaries between these categories cannot be established in real time.

We include all of these (the discrepancies, the interpolations, the unverified records) because removing them would produce a dataset that is more wrong, not less. It would present a false consensus where none exists. A dataset that only contains verified facts about an ongoing conflict is either very small or very dishonest. We have chosen to be large and transparent. Each monthly release will carry forward this commitment: new discrepancies will be documented, not silently resolved.

6.3 The Training Loop and the Synthetic Turn

This dataset will not remain in a GitHub repository. It will be indexed by web crawlers, summarized by AI systems, cited in analyses that are themselves crawled and summarized, and ingested into the training corpora of future large language models.³ This is not a speculative concern. It is the mechanical reality of how knowledge circulates in the current information ecosystem. And with each monthly release, the dataset's footprint in that ecosystem grows.

The implications are significant. A model trained on a corpus that includes this dataset (or, more likely, trained on content that summarized or referenced this dataset) will carry forward our interpretive judgments, our confidence ratings, our interpolations, and our errors. It will do so without the provenance markers we have attached. The model will not know that a particular casualty figure was rated MEDIUM confidence, or that a war cost estimate was interpolated from two anchor points. It will know only that the number was X.⁴

This is the ouroboros at the center of the synthetic turn in AI development. Models are increasingly trained on data generated or curated by other models (Bender et al., 2021; Bommasani et al., 2021). Summary layers reference summary layers. Confidence intervals are stripped as data moves downstream.⁵ The original source, be it a Reuters wire report, a CENTCOM press release, a Red Crescent field count, becomes archaeologically distant from the synthetic output it eventually informs. Recent empirical work confirms the danger: Shumailov et al. (2024) demonstrate that models trained recursively on model-generated data suffer irreversible distributional collapse, losing the tails of the original data distribution. Related work on language-model hallucination (Ji et al., 2023) shows that confident but ungrounded generation is not an edge case but a persistent property of current systems.

We do not claim to solve this problem. We claim only that naming it is a precondition for addressing it, and that datasets which are transparent about their own synthetic and interpretive elements contribute more honestly to the information ecosystem than those which are not. If a

future model ingests this dataset, or any of its monthly successors, it will at least have the opportunity to ingest the caveat alongside the claim.

6.4 The Flood and the Counter-Technology: Structural and Strategic Layers

Earlier sections referred to the information flood in two layers. We now elaborate upon those layers because they require different responses.

The structural layer. The structural information flood is the condition in which the rate of information production, whether organic or synthetic, exceeds the rate at which individuals and institutions can evaluate what is true, what is uncertain, what is fabricated, and what is algorithmically generated. This framing is consistent with the public-health literature on infodemics, which treats the problem not only as falsehood but also as an overabundance of accurate, inaccurate, contradictory, or incomplete information that can impede timely action (Borges do Nascimento et al., 2022; Chiou et al., 2022; World Health Organization, 2021). It is an epistemic crisis, not merely a volumetric one. The problem is not that we have too many data points about the 2026 US–Iran conflict. The problem is that we lack shared infrastructure for evaluating which of those data points to trust. This layer is largely emergent: no single actor designed it; it arises from the aggregate behavior of platforms, summarization systems, citation chains, and downstream reuse practices that strip provenance as content moves.

The strategic layer. The strategic information flood is what happens when actors deliberately exploit the structural layer for political or economic ends. The political-science literature on information operations treats disinformation not as noise but as a strategic instrument (Rid, 2020). States, militaries, and non-state actors invest significant resources in shaping what gets reported, how events are framed, and what interpretations circulate, not as a byproduct of conflict but as a deliberate extension of it. The 2026 US–Iran conflict has generated documented evidence of strategic flooding: coordinated AI-generated content, synchronized posting bursts across platforms, and multilingual accounts targeting geographically distant audiences, consistent with centrally coordinated influence rather than organic response (Cyabra, 2026; Murray & Linvill, 2026). This is not the flood passively rising. It is actors opening the floodgates.

The conventional response to the structural layer is curation: gatekeeping, editorial selection, and expert consensus. These mechanisms remain valuable, but they operate at timescales that are often inadequate to real-time conflicts. By the time retrospective datasets release fully verified event data for this conflict, the political and analytical window for many research questions will have narrowed or closed.

We propose that a different kind of response is also needed: datasets that carry their own epistemic infrastructure with them. Confidence scores. Source attribution. Version histories.

Explicit documentation of what was interpolated, what was unverified, what was generated by an AI system, and under what instructions. These features do not make the data true. They make the data assessable. They restore to downstream consumers some of the judgment capacity that the flood otherwise strips away.

This is what we mean by a counter-technology: not a technology that opposes the flood by restricting information, but one that opposes it by restoring the conditions of evaluation. The IranWar.ai dataset is an experiment in this mode: structured data, released at speed, updated through daily dashboard data and scheduled versioned research releases, with its uncertainty visible. Whether this transparency survives contact with the information ecosystem (i.e., whether confidence scores and caveats propagate downstream or are stripped in summarization) is itself an empirical question. We invite the research community to help us answer it.

Crucially, transparency operates differently against each layer of the flood. Against the structural layer, transparency is both necessary and often sufficient: provenance, confidence scores, and version history give downstream consumers what they need to recover accessibility. Against the strategic layer, transparency remains necessary, but it is not sufficient. A coordinated influence operation that produces confidently wrong content at scale is not corrected by labeling our data more carefully. It is corrected only when interpretive expertise meets the labels, and that requires Section 6.5.

6.5 The Information Environment as a Political Variable

Section 6.4 names the strategic layer of the information flood; this section develops what it implies for analysis. Once we accept that the information environment is something actors deliberately produce, the limitations of the dataset begin to look different. They cease to be uniformly "data quality" issues and become, at least in part, observable traces of political behavior.

The dataset itself contains evidence of this. The divergence between Iranian Red Crescent casualty figures and those reported by HRANA, the diaspora human rights organization, is not primarily a data-quality problem amenable to better collection methods. It is a political fact: the direct expression of two actors with irreconcilable interests in what the casualty numbers say. The Iranian state has structural incentives to minimize civilian-casualty figures to limit international pressure; diaspora human rights organizations have structural incentives to maximize them to delegitimize the regime. The discrepancy between their figures is not noise to be resolved but signal to be interpreted.

Gibilisco and Steinberg's (2023) formal model of strategic reporting in conflict data, which this paper has cited in earlier sections, speaks directly to this dynamic. Their core insight (that governments strategically time and frame their reporting of conflict events to manipulate the

inferences of outside observers) applies not only to the sources the dataset draws on but to the dataset's own interpretive choices. Every decision about which source to weight, how to code a contested event, and what confidence level to assign is an interpretive act embedded in a political context. Acknowledging this does not undermine the dataset's credibility. On the contrary, the acknowledgment is a precondition for using it responsibly.

Practically, this argues for a working stance we recommend to all downstream users of the dataset: read discrepancies as political facts before reading them as errors. When two source categories disagree, the disagreement is itself an observation. Researchers should disaggregate by source between US government, Iranian government, ACLED, wire services, diaspora monitors, and treat the variance as a measurement of contestation rather than as a problem to control away. The dataset's record-level source attribution is designed to support exactly this. Used in this way, the limitations catalogued in Section 5 become inputs to *political* analysis as much as constraints on *statistical* analysis.

This refinement leaves the counter-technology framework intact but sharpens it. Transparent datasets are an effective counter-technology against the structural layer of the information flood. Against the strategic layer, transparency alone is insufficient; it must be paired with *interpretive expertise*: the capacity to read discrepancies as political facts and intentional choices, and to situate data within the actor incentives and institutional contexts that produced it. This is precisely the contribution that political science, with its frameworks for analyzing strategic behavior, institutional interests, and information asymmetries, is positioned to make. We hope this dataset becomes a useful instrument for that work.

7 Research Applications

The monthly release cadence was designed with longitudinal research in mind. Many of the most interesting questions about this conflict (i.e., how did escalation dynamics shift? At what points did financial markets price in the military risk?) can only be answered with data that grows over time. What follows is not exhaustive, but offers starting points across disciplines.

7.1 Political Science and International Relations

The most immediate applications are in escalation analysis: daily strike counts, shifts in target category over time, and the geographic spread of strikes across provinces can support both qualitative case studies and quantitative sequence analysis. Retaliation patterns across the 307 retaliation events offer a basis for deterrence research. The 66 diplomatic events, cross-referenced against the military tempo data, create a natural dataset for studying how negotiation dynamics interact with battlefield activity.

The IranWar.ai dataset offers unusually rich material for political science research precisely because it captures an active, high-tempo, high-stakes conflict in near real time: a data environment that existing conflict datasets like UCDP and ACLED, with their weeks-to-months reporting lags, cannot provide. Three theoretical frameworks from the IR literature map directly onto the dataset's structure and suggest immediate research applications.

Escalation dynamics and coercive logic. The dataset's 641 strike events and 307 retaliation events, sequenced by day across the 30-day conflict event window, create a natural foundation for testing competing theories of coercive bargaining. Schelling's (1966) foundational distinction between compellence (coercive action designed to force a new action or stop an ongoing action) and deterrence (coercive action designed to prevent an action) generates a testable question from this data: does the US/Israeli strike tempo follow a compellence logic, escalating to force Iranian concessions on the Strait of Hormuz, or a punishment logic, degrading Iranian military capacity regardless of behavioral response? The day-level sequencing of strikes, retaliations, and diplomatic events allows researchers to examine whether escalatory moves preceded, followed, or were independent of negotiation attempts, and whether the answer differs from the US/Israeli and Iranian perspectives. That question sits at the center of the bargaining model of war, where observable battlefield behavior and public diplomatic signals become evidence about information asymmetries, commitment problems, and the credibility of coercive threats (Fearon, 1995).

Credible commitment and public signaling. Fearon's (1995) rationalist framework for war emphasizes the role of information asymmetries and credible commitment problems in explaining why states fight rather than negotiate. The 2026 US–Iran conflict presents an unusually legible information environment for testing these dynamics: both sides have made detailed public claims about what they are striking and why, generating a rich record of official signaling that the dataset captures through CENTCOM press releases, IDF statements, and Iranian government sources. This allows researchers to examine whether public signaling tracks with actual targeting patterns (whether states are doing what they say they are doing) and to assess the credibility of commitment claims against observable battlefield behavior. The dataset's confidence scoring and source attribution make it possible to distinguish between events that are well-verified and those that remain contested, which is itself analytically relevant: in Fearon's framework, the credibility of a signal depends partly on the cost of making it, and the gap between claimed and verifiable action is a direct measure of that credibility.

The conflict also offers a rare opportunity for cross-regime comparison of audience cost dynamics (Fearon, 1994). Trump's public ultimatums on Truth Social (some of which contradicted prior US positions), Netanyahu's maximalist public declarations about Iranian nuclear facilities, and Iranian leaders' statements to a domestic audience that cannot punish them electorally but can

delegitimize them through religious and nationalist frameworks represent three distinct audience-cost structures operating simultaneously. The dataset's capture of official statements alongside observable battlefield behavior allows researchers to test whether the credibility and follow-through of public commitments varied systematically across these three very different domestic political contexts.

Conflict diffusion across actors and fronts. The dataset's multi-domain structure, which tracks military, diplomatic, financial, humanitarian, and naval events simultaneously, enables research on how conflict spreads across actors and issue areas over time. The secondary fronts in Lebanon and Yemen are currently coded as contextual data but represent analytically important cases of conflict diffusion: points at which the bilateral US–Iran dynamic expanded to involve Hezbollah and Houthi actors with distinct interests, capabilities, and relationships to the primary belligerents. Building on the diffusion literature (Buhaug & Gleditsch, 2008), researchers can use the dataset's daily event sequencing to examine the triggers and timing of these expansions, and to ask whether diffusion followed military, diplomatic, or economic escalation signals from the primary theater. Stated concretely: why did the conflict diffuse to Lebanon and Yemen, and not with equal intensity to other plausible secondary theaters such as Syria or Iraq?

7.2 Economics and Financial Markets

The financial domain (301 events) integrates oil price movements with the strike and retaliation timeline, enabling event-study designs that link specific military events to commodity price responses. The market-sector data—tracking defense, airline, and broader equity indices from a pre-conflict baseline—supports portfolio and sectoral analysis. Tanker transit records provide a proxy for the economic effects of the Strait of Hormuz disruption. We openly invite economists and financial analysts to broaden the related research agenda accessible to this dataset.

7.3 Public Health, Epidemiology, and Civilian Harm Exposure

The humanitarian domain (173 events) includes daily casualty estimates disaggregated by faction, geographic information on infrastructure attacks, and specific events such as the Minab school strike that drew international attention. These records support population-level exposure analysis, civilian casualty pattern analysis, and healthcare system degradation research. Researchers in this area should read Section 5 carefully regarding the limitations of the casualty figures.

Beyond descriptive uses of the humanitarian fields, the dataset's cross-domain structure supports a more substantive contribution to conflict public health: the construction of a regional, time-varying composite index of civilian harm exposure. We outline that index here as a methodological proposal that the dataset is designed to enable, rather than as a finished instrument;

full operationalization, validation, and sensitivity analysis will be presented in a separate companion methodology document.

7.3.1 The Population Harm Exposure Index (PHEI): Rationale

Civilian harm in armed conflict cannot be captured by mortality alone. Deaths are the most visible outcome, but they are also the most under-reported, the most politically contested, and the most lagged relative to the events that cause them (Checchi & Roberts, 2008; Roberts et al., 2010; Jawad et al., 2020). A growing body of work in conflict public health argues for composite measures that combine kinetic exposure, infrastructure damage, displacement, and health-system stress into population-level indicators that can be tracked at sub-national resolution and over time (Salama, Spiegel, Talley, & Waldman, 2004; Heudtlass, Speybroeck, & Guha-Sapir, 2016). Composite indices of this kind are now widely used in humanitarian operations: the INFORM Severity Index aggregates more than 30 indicators across three dimensions (impact, conditions of affected people, and complexity) into a globally comparable severity score (Poljanšek, Marin Ferrer, & Disperati, 2020), and ACLED's Conflict Exposure measure couples event data with population grids to estimate populations within 1, 2, and 5 km of each event (Raleigh, Linke, Hegre, & Karlsen, 2010; ACLED, 2023).

We deliberately use the term harm exposure rather than trauma to keep the construct firmly within an epidemiological exposure-assessment frame and to avoid the clinical and psychometric commitments that come with trauma-specific instruments. The PHEI is a population-level indicator of environmental exposure to conflict-related harm at a defined geographic and temporal resolution. It is conceptually closer to environmental-exposure indices (e.g., particulate matter dose-response work) than to individual-level trauma scales, and it is designed to be computable from open-source event data without individual-level survey instruments. This framing is complementary to, rather than a substitute for, the trauma-psychology lens developed elsewhere: the PHEI captures population-level exposure conditions that shape the prevalence of the clinical trauma outcomes that other instruments measure directly.

7.3.2 Proposed Domains and Indicators

The PHEI aggregates indicators across five domains, each computable from variables already present in the IranWar.ai dataset or in routine humanitarian data feeds. We follow the construction logic recommended for composite indicators in humanitarian contexts (Poljanšek, Marin Ferrer, & Disperati, 2020) and the dose-response framing common in environmental epidemiology (Salama et al., 2004).

- 1. Kinetic intensity.** Strike and retaliation event counts within a defined geographic unit (province or 5-km grid cell, following the ACLED proximity convention) and a defined temporal window (rolling 7-day or 30-day). Where weapon-class metadata is available,

events are weighted by likely lethal radius. This domain captures direct exposure to violent events.

2. **Infrastructure damage.** Counts and severity codes for events affecting civilian-facing infrastructure, with a particular weight on healthcare, water, energy, and education facilities. Damage to healthcare infrastructure, in particular, is a strong and well-documented predictor of indirect mortality from non-battle causes in conflict-affected populations (Jawad et al., 2020).
3. **Casualty estimates.** Daily killed-and-injured estimates by faction, drawn from the humanitarian domain and tagged with their data_confidence rating so the index does not silently absorb low-confidence figures. Following Checchi and Roberts (2008), we recommend reporting raw and confidence-weighted variants in parallel rather than collapsing them.
4. **Displacement and access disruption.** Where available from UNHCR, OCHA, IOM, and Red Crescent feeds: estimated displaced population, population in areas with internet or telecommunications blackout, and population in areas with restricted humanitarian access. Excess mortality in refugee and IDP populations during complex emergencies is well established (Heudtlass, Speybroeck, & Guha-Sapir, 2016) and is poorly captured by event counts alone.
5. **Geographic clustering and persistence.** A spatial-temporal autocorrelation term that distinguishes a single isolated event from a sustained campaign in the same locality. This domain operationalizes the intuition that repeated exposure in the same neighborhood compounds harm in ways that a simple sum across the period does not capture.

7.3.3 Construction and Calibration

Within each domain, indicators are normalized (typically min-max within the period of analysis), aggregated to a domain sub-score using either an arithmetic mean or a geometric mean depending on substitutability, and then combined across domains using documented weights. The INFORM Severity methodology (Poljanšek et al., 2020) and validation work on composite humanitarian indices (Lopez et al., 2023) provide established templates: exploratory and confirmatory factor analysis can be used to test whether the proposed domains load on a single latent construct of "crisis severity" or "harm exposure," and to refine indicator selection. We anticipate four severity bands (LOW / MODERATE / HIGH / EXTREME) set on the basis of percentile cuts within the in-conflict distribution rather than on absolute thresholds, both because absolute thresholds in this conflict are not yet well calibrated and because relative ranking is what matters for resource-allocation decisions.

7.3.4 Intended Uses and Limits

The PHEI is intended to support, not replace, existing assessment instruments. Three intended uses are immediate. First, region-level severity scores allow rapid comparison of harm intensity across Iran's provinces and across the secondary fronts in Lebanon and Yemen. Second, the time-varying score supports surveillance-style analysis of whether harm exposure is rising or falling within a given region, independently of national casualty totals. Third, the index can serve as an exposure variable in panel-style models of indirect mortality of the kind developed by Jawad and colleagues (2020), with the country-month design replaced by province-week or grid-week observations.

The limitations of the PHEI follow from the limitations of the underlying data and are honest by design. Casualty inputs inherit the discrepancies catalogued in Section 5 and Section 6. Infrastructure-damage codes depend on satellite-imagery analysis and humanitarian reporting that lag the events they describe. Displacement figures from Iran are particularly sparse during internet blackouts. We therefore recommend reporting the PHEI with explicit confidence intervals derived from the data_confidence ratings of its inputs, and treating low-confidence regional scores as hypotheses to be revised in the next monthly release rather than as findings.

A separate methodology paper, prepared in parallel with future monthly releases, will document the index construction in full, including the indicator-by-indicator calibration, sensitivity to weighting choices, and validation against independent humanitarian assessments where available. The PHEI is meant to be auditable, contestable, and revisable: properties we treat as features rather than concessions, in keeping with the broader counter-technological commitments of this paper.

7.4 Military Science

The strike dataset (641 events across 61 locations in 26 provinces) supports air campaign analysis, including target selection patterns, tempo variation, and comparison with historical precedents such as Iraq 2003, the Gulf War, and the Libya campaign. Naval data (28 events) covers carrier deployments, vessel incidents, and Hormuz-related events relevant to blockade and maritime strategy research. We openly invite military scholars to broaden the related research agenda accessible to this dataset.

7.5 Media, Communication, and Information Science

The event_description field across 1,653 rows represents a substantial corpus of conflict framing from a consistently structured source, with corresponding source attribution metadata. This creates opportunities for framing analysis and reporting-lag studies. More broadly, the dataset's structure (i.e., AI-assembled, provenance-tagged, released with known uncertainty)

makes it an object of study in its own right for researchers working on synthetic data lineage, model collapse, and epistemic infrastructure in AI systems. Studies that combine our records with the documented Iranian AI-driven influence operations reported by Cyabra (2026) and Murray and Linvill (2026) are particularly well-positioned to model the strategic layer of the information flood empirically. We openly invite researchers of the sociological, media, journalistic, and information sciences to broaden the related research agenda accessible to this dataset.

8 Invitation to Collaborate

We are releasing this dataset comprising the first 30 days of conflict, which means it is incomplete in ways we can already identify and probably in ways we cannot yet see. The collaboration invitation is therefore genuine, not formal.

8.1 Data Validation

The most valuable contributions would come from researchers with regional expertise in Iranian studies, Gulf politics, Israeli security, or Houthi and Hezbollah operations who can cross-reference our records against local-language sources we have not accessed. Four retaliation records are currently unverified; any of the casualty discrepancies documented in Section 5 would benefit from expert scrutiny. Validated corrections will be incorporated into the next monthly release with full attribution.

8.2 Data Extension

The most obvious gaps in the current version are sub-national granularity (province-level coding exists, but district-level coding for Iranian strike locations would improve spatial analysis), casualty disaggregation (current composite figures could be broken into more structured fields with the right expertise), and economic variables (exchange-rate data, shipping insurance indices, and detailed commodity prices would significantly enrich the financial domain). Two additional extensions deserve specific mention because contributors have already raised them and because they sit at the edges of the current scope.

Displacement and migration flows. The current taxonomy treats displacement as a humanitarian variable but does not provide a dedicated migration-flow layer. Linking our data to UNHCR, IOM, and ACLED population-on-the-move feeds would produce a more direct connection between kinetic events and population displacement, and would support a more developed PHEI displacement domain (Section 7.3). We see this as a natural target for a future release.

Civilian observation reporting. Future versions of the project may incorporate structured anonymous observation reports from civilians in affected areas (e.g., what was seen, where, and when), submitted through a dedicated reporting form rather than through open file uploads. Such reports would not be treated as verified facts. They would enter the pipeline as additional candidate records, tagged LOW confidence by default, and triangulated against media, NGO, and satellite sources before being upgraded. This approach would broaden the diversity of inputs without compromising the dataset's confidence-scoring discipline. We treat this as an explicit research and ethics question (covering consent, identification risk, and the political risks to reporters in conflict zones) and will document the design decisions for any future implementation in the methodology companion.

8.3 Methodological Contributions

Some events likely appear in both the *timeline-events.json* file and the strike files, which creates potential double-counting. Systematic deduplication methods would be a meaningful contribution. Confidence calibration (i.e., empirically assessing our High/Medium/Low ratings against subsequently confirmed reporting) would help future users weigh the data appropriately.

8.4 Analysis and Publication

We encourage researchers to use this dataset for peer-reviewed publications. We ask only for an appropriate citation and that any corrections or extensions be contributed back to the repository.

How to Contribute. Fork the repository at github.com/jethomasphd/WarTheater, add your data, corrections, or analysis to the ResearchData/ directory, and submit a pull request with a clear description of what changed and why. For coordination or questions, open an issue on the repository.

9 Ethics Statement

This dataset documents a real, ongoing armed conflict. The data includes records of civilian casualties, attacks on schools and hospitals, and other events involving human suffering. Researchers using this data should:

1. Treat casualty figures with appropriate gravity and uncertainty. These numbers represent real human lives.
2. Acknowledge the limitations of OSINT data in conflict zones, particularly regarding civilian harm.
3. Be transparent about the provenance and limitations of any findings.

4. Consider the potential for their work to be misused or misinterpreted.
5. Follow their institutional IRB/ethics board guidance for research involving conflict data.
6. **Propagate uncertainty.** Researchers who cite specific figures from this dataset (e.g., casualty counts, cost estimates, strike tallies) should carry forward the associated confidence ratings and source attributions. Stripping uncertainty metadata from data points and presenting them as established facts is a form of information loss that this dataset is specifically designed to resist.

The authors affirm that this dataset was compiled exclusively from publicly available open-source intelligence. No classified, proprietary, or restricted-access data was used.

10 Conclusion

The IranWar.ai Event-Level Research Dataset provides the research community with immediate, structured access to multi-domain data from the 2026 US–Iran conflict. Its 1,653 event-level observations across military, financial, humanitarian, and diplomatic domains with full source attribution, confidence scoring, and reproducible extraction code lowers the barrier to rigorous quantitative analysis of this conflict.

But the dataset is also something else: a test case for how structured data behaves in the information flood. It is assembled in part by AI systems. It contains known imperfections. It will be crawled, summarized, and almost certainly ingested into future AI training corpora. In releasing it with its uncertainty visible (i.e., confidence scores intact, limitations cataloged, interpretive chain documented) we are conducting a small experiment in epistemic transparency.

The flood we have described is dual-layered. Its structural layer is largely a byproduct of how information now circulates; its strategic layer is something actors make on purpose. Transparent datasets are a meaningful counter-technology against the structural layer. Against the strategic layer, transparency is necessary but insufficient and must be paired with interpretive expertise. The IranWar.ai dataset is built to support both responses: provenance and confidence scoring for the structural layer, and source-disaggregated records that make political contestation legible for the strategic layer.

The question this experiment asks is simple, and it applies far beyond this dataset: Does it matter? Does labeling uncertainty change how data propagates? Does provenance survive summarization? Does the downstream consumer (human or machine) preserve the caveat, or only the claim? We do not know these answers. But we believe these questions are among the most important facing the research community in an era when the line between observed data, curated data, and synthetic data is dissolving. Whether or not the research community has chosen to name

it, this is the environment in which evidence now circulates. The infrastructure we build now (i.e., the habits of transparency, the norms of provenance, the willingness to release imperfect data alongside its own critique) will determine whether the next generation of knowledge systems inherits our evidence or only our confidence.

The dataset will be released on a scheduled monthly cadence for the duration of the conflict and its aftermath, with each research release versioned and snapshotted. Interim correction releases may be issued when source updates, major errors, or analytically important reclassifications warrant earlier publication. We invite researchers across disciplines to validate, extend, and analyze this data, and to contribute their findings and improvements back to the open-source repository.

Data Availability Statement

The dataset, codebook, extraction script, source data, and documentation are available at:

Repository: <https://github.com/jethomasphd/WarTheater>

Dataset directory: ResearchData/

Live dashboard: <https://iranwar.ai>

Release cadence: Daily live dashboard data; monthly, versioned research releases (v1.0, v1.1, ...), with interim correction releases when warranted.

Acknowledgments

The IranWar.ai dashboard and dataset rely on the work of journalists, conflict monitors, humanitarian organizations, and government transparency offices whose reporting makes open-source intelligence possible. We particularly acknowledge the Armed Conflict Location & Event Data Project (ACLED), the Critical Threats Project at AEI, the USNI Fleet Tracker, and the IAEA for their systematic public reporting.

References

- ACLED. (2023). Conflict exposure methodology. Armed Conflict Location & Event Data Project. <https://acleddata.com/methodology/conflict-exposure-methodology>
- Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the dangers of stochastic parrots: Can language models be too big? Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT '21), 610–623. <https://doi.org/10.1145/3442188.3445922>

- Borges do Nascimento, I. J., Pizarro, A. B., Almeida, J. M., Azzopardi-Muscat, N., Gonçalves, M. A., Björklund, M., & Novillo-Ortiz, D. (2022). Infodemics and health misinformation: A systematic review of reviews. *Bulletin of the World Health Organization*, 100(9), 544–561. <https://doi.org/10.2471/BLT.21.287654>
- Bommasani, R., Hudson, D. A., Adeli, E., et al. (2021). On the opportunities and risks of foundation models. arXiv preprint, arXiv:2108.07258. <https://doi.org/10.48550/arXiv.2108.07258>
- Buhaug, H., & Gleditsch, K. S. (2008). Contagion or confusion? Why conflicts cluster in space. *International Studies Quarterly*, 52(2), 215–233. <https://doi.org/10.1111/j.1468-2478.2008.00499.x>
- Checchi, F., & Roberts, L. (2008). Documenting mortality in crises: what keeps us from doing better. *PLoS medicine*, 5(7), e146. <https://doi.org/10.1371/journal.pmed.0050146>
- Chiou, H., Voegeli, C., Wilhelm, E., Kolis, J., Brookmeyer, K., & Prybylski, D. (2022). The future of infodemic surveillance as public health surveillance. *Emerging Infectious Diseases*, 28(13), S121–S128. <https://doi.org/10.3201/eid2813.220696>
- Cyabra. (2026, March 13). Iran's coordinated campaign: Deepfakes used to manipulate the war narrative. <https://cyabra.com/reports/irans-coordinated-campaign-deepfakes-used-to-manipulate-the-war-narrative/>
- Eck, K. (2012). In data we trust? A comparison of UCDP GED and ACLED conflict events datasets. *Cooperation and Conflict*, 47(1), 124–141. <https://doi.org/10.1177/0010836711434463>
- Fearon, J. D. (1994). Domestic political audiences and the escalation of international disputes. *The American Political Science Review*, 88(3), 577–592. <http://www.jstor.org/stable/2944796>
- Fearon, J. D. (1995). Rationalist explanations for war. *International Organization*, 49(3), 379–414. <http://www.jstor.org/stable/2706903>
- Geburu, T., Morgenstern, J., Vecchione, B., Vaughan, J. W., Wallach, H., Daumé III, H., & Crawford, K. (2021). Datasheets for datasets. *Communications of the ACM*, 64(12), 86–92. <https://doi.org/10.48550/arXiv.1803.09010>
- Gibilisco, M., & Steinberg, J. (2023). Strategic reporting: A formal model of biases in conflict data. *American Political Science Review*, 117*(4), 1291–1307. <https://doi.org/10.1017/S0003055422001162>
- Gleditsch, N. P., Wallensteen, P., Eriksson, M., Sollenberg, M., & Strand, H. (2002). Armed Conflict 1946–2001: A New Dataset. *Journal of Peace Research*, 39(5), 615–637. <http://www.jstor.org/stable/1555346>

- Hegre, H., Croicu, M., Eck, K., & Högbladh, S. (2020). Introducing the UCDP Candidate Events Dataset. *Research & Politics*, 7(3). <https://doi.org/10.1177/2053168020935257>
- Heudtlass, P., Speybroeck, N., & Guha-Sapir, D. (2016). Excess mortality in refugees, internally displaced persons and resident populations in complex humanitarian emergencies (1998-2012) - insights from operational data. *Conflict and health*, 10, 15. <https://doi.org/10.1186/s13031-016-0082-9>
- Jawad, M., Hone, T., Vamos, E. P., Roderick, P., Sullivan, R., & Millett, C. (2020). Estimating indirect mortality impacts of armed conflict in civilian populations: panel regression analyses of 193 countries, 1990-2017. *BMC medicine*, 18(1), 266. <https://doi.org/10.1186/s12916-020-01708-5>
- Ji, Z., Lee, N., Frieske, R., et al. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), 1–38. <https://doi.org/10.1145/3571730>
- Lopez, V. K., Nika, A., Blanton, C., Talley, L., & Garfield, R. (2023). Can severity of a humanitarian crisis be quantified? Assessment of the INFORM severity index. *Globalization and health*, 19(1), 7. <https://doi.org/10.1186/s12992-023-00907-y>
- Miller, E., Kishi, R., Raleigh, C., & Dowd, C. (2022). An agenda for addressing bias in conflict data. *Scientific data*, 9(1), 593. <https://doi.org/10.1038/s41597-022-01705-8>
- Murray, E., & Linvill, D. (2026). From Texas to Tehran: A multilingual, IRGC-affiliated influence operation on X, Instagram, and Bluesky. *Clemson University Media Forensics Hub Reports*. https://open.clemson.edu/mfh_reports/10/
- Palmer, G., McManus, R. W., D’orazio, V., Kenwick, M. R., Karstens, M., Bloch, C., ... & Soules, M. J. (2022). The MID5 dataset, 2011–2014: Procedures, coding rules, and description. *Conflict Management and Peace Science*, 39(4), 470-482. <https://ideas.repec.org/a/sae/compssc/v39y2022i4p470-482.html>
- Paullada, A., Raji, I. D., Bender, E. M., Denton, E., & Hanna, A. (2021). Data and its (dis)contents: A survey of dataset development and use in machine learning research. *Patterns* (New York, N.Y.), 2(11), 100336. <https://doi.org/10.1016/j.patter.2021.100336>
- Poljanšek, K., Marin Ferrer, M., & Disperati, S. (2020). INFORM Severity Index: Concept and methodology. Publications Office of the European Union, JRC Technical Report. <https://drmkc.jrc.ec.europa.eu/inform-index/INFORM-Severity>
- Raleigh, C., Linke, A., Hegre, H., & Karlsen, J. (2010). Introducing ACLED: An Armed Conflict Location and Event Dataset. *Journal of Peace Research*, 47(5), 651–660. <http://www.jstor.org/stable/20798933>
- Rid, T. (2020). *Active measures: The secret history of disinformation and political warfare*. Farrar, Straus and Giroux.

- Roberts, B., Morgan, O. W., Sultani, M. G., Nyasulu, P., Rwebangila, S., Myatt, M., Sondorp, E., Chandramohan, D., & Checchi, F. (2010). A new method to estimate mortality in crisis-affected and resource-poor settings: validation study. *International journal of epidemiology*, 39(6), 1584–1596. <https://doi.org/10.1093/ije/dyq188>
- Salama, P., Spiegel, P., Talley, L., & Waldman, R. (2004). Lessons learned from complex emergencies over past decade. *Lancet* (London, England), 364(9447), 1801–1813. [https://doi.org/10.1016/S0140-6736\(04\)17405-9](https://doi.org/10.1016/S0140-6736(04)17405-9)
- Schelling, T. C. (1966). *Arms and influence*. Yale University Press. <https://yalebooks.yale.edu/book/9780300246742/arms-and-influence/>
- Shumailov, I., Shumaylov, Z., Zhao, Y., Papernot, N., Anderson, R., & Gal, Y. (2024). AI models collapse when trained on recursively generated data. *Nature*, 631(8022), 755–759. <https://doi.org/10.1038/s41586-024-07566-y>
- Sundberg, R., & Melander, E. (2013). Introducing the UCDP Georeferenced Event Dataset. *Journal of Peace Research*, 50(4), 523–532. <https://doi.org/10.1177/0022343313484347>
- Weidmann, N. B. (2015). On the accuracy of media-based conflict event data. *Journal of Conflict Resolution*, 59(6), 1129–1149. <https://ideas.repec.org/a/sae/jocore/v59y2015i6p1129-1149.html>
- World Health Organization. (2021). *Infodemic management: An overview of infodemic management during COVID-19, January 2020–May 2021*. <https://www.who.int/publications/i/item/9789240035966>

Notes

- ¹ The term *information flood* is used here in a specific technical sense: a dual-layer epistemic crisis comprising both a structural component (volume, velocity, and synthetic origin overwhelming evaluation infrastructure) and a strategic component (deliberate adversarial production of disinformation by state and non-state actors). The structural framing draws on epidemiological models of environmental exposure, treating information not as neutral content but as an ambient condition with dose-response characteristics (World Health Organization, 2021). The strategic framing draws on the political-science literature on information operations (e.g., Rid, 2020).
- ² We use "counter-technology" to describe any system designed not to accelerate information production but to restore the conditions under which information can be meaningfully evaluated: provenance, uncertainty quantification, and interpretive context. Also, see Gebru et al. (2021) on datasheets as documentation infrastructure for datasets.
- ³ This is not hypothetical. Common Crawl, the primary web corpus used to train large language models (Bommasani et al., 2021), indexes static sites like IranWar.ai. Any content published to the open web is a candidate for inclusion in future training corpora.
- ⁴ The phenomenon of models trained on model-generated data has been termed "model collapse" in the machine learning literature. Shumailov et al. (2024) demonstrate that indiscriminate use of model-generated content in training causes irreversible defects in resulting models, in which tails of the original content distribution disappear.
- ⁵ Paullada et al. (2021) document a broader pattern in which the provenance, composition, and limitations of training datasets are systematically underexamined in machine learning research—a problem that compounds as datasets are reused across contexts.